

OBJECTIVE OF ANALYSIS

The major objective of the analysis is to visualize salaries based on different metrics like gender, experience, and education for the IT industry in United States of America from the Survey dataset.

RESEARCH QUESTION

Are male software engineers paid higher salaries than their female and non-binary counterparts ?

MOTIVATION FOR CHOOSING THIS DATASET

This dataset consists of various types of inputs and inconsistent inputs that makes it ideal for data cleaning practice.

IMPORTING NECESSARY LIBRARIES

```
In [ ]: import pandas as pd
import numpy as np
import difflib
import fuzzywuzzy
from fuzzywuzzy import process
import seaborn as sns
import matplotlib.pyplot as plt
import chardet
import openpyxl

/Users/anshumangupta/Desktop/portfolio/Data Cleaning/Data Cleaning/salary/lib/python3.10/site-packages/fuzzywuzzy/fuzz.py:1
1: UserWarning: Using slow pure-python SequenceMatcher. Install python-Levenshtein to remove this warning
warnings.warn('Using slow pure-python SequenceMatcher. Install python-Levenshtein to remove this warning')
```

IMPORTING DATA (.csv)

```
In [ ]: # Loading raw data into pandas dataframe
raw_data = pd.DataFrame(pd.read_csv('data/survey.csv'))

# Viewing first 5 rows
raw_data.head(5)
```

Out[]:

	Timestamp	How old are you?	What industry do you work in?	Job title	If your job title needs additional context, please clarify here:	What is your annual salary? (You'll indicate the currency in a later question. If you are part-time or hourly, please enter an annualized equivalent -- what you would earn if you worked the job 40 hours a week, 52 weeks a year.)	How much additional monetary compensation do you get, if any (for example, bonuses or overtime in an average year)? Please only include monetary compensation here, not the value of benefits.	Please indicate the currency	If "Other," please indicate the currency here:	If your income needs additional context, please provide it here:	What country do you work in?	If you're in the U.S., what state do you work in?	What is your gender?
0	4/27/2021 11:02:10	25-34	Education (Higher Education)	Research and Instruction Librarian	NaN	55,000	0.0	USD	NaN	NaN	United States	Massachusetts	White
1	4/27/2021 11:02:22	25-34	Computing or Tech	Change & Internal Communications Manager	NaN	54,600	4000.0	GBP	NaN	NaN	United Kingdom	NaN	Canadian
2	4/27/2021 11:02:38	25-34	Accounting, Banking & Finance	Marketing Specialist	NaN	34,000	NaN	USD	NaN	NaN	US	Tennessee	Chattanooga
3	4/27/2021 11:02:41	25-34	Nonprofits	Program Manager	NaN	62,000	3000.0	USD	NaN	NaN	USA	Wisconsin	Millwaukee
4	4/27/2021 11:02:42	25-34	Accounting, Banking & Finance	Accounting Manager	NaN	60,000	7000.0	USD	NaN	NaN	US	South Carolina	Greenville

CHECKING NUMBER OF ROWS AND COLUMNS IN THE DATASET

```
In [ ]: # Checking number of rows and columns
raw_data.shape
```

```
Out[ ]: (27893, 18)
```

CHANGING COLUMN HEADERS FOR SIMPLICITY

```
In [ ]: # Setting new column headers
new_column_headers = ['timestamp', 'age_group', 'curr_industry', 'job_title', 'job_context', 'curr_salary', 'compensation', 'currency', 'additional_currency', 'income_context', 'country']
init_data = raw_data.set_axis(new_column_headers, axis=1, inplace=False)
init_data.head(5)
```

/var/folders/h2/3d918ncd5xdgjlh3ffmkqyk0000gn/T/ipykernel_39323/4246118180.py:3: FutureWarning: DataFrame.set_axis 'inplace' keyword is deprecated and will be removed in a future version. Use `obj = obj.set_axis(..., copy=False)` instead

```
init_data = raw_data.set_axis(new_column_headers, axis=1, inplace=False)
```

```
Out[ ]: timestamp age_group curr_industry job_title job_context curr_salary compensation currency additional_currency income_context country
```

0	4/27/2021 11:02:10	25-34	Education (Higher Education)	Research and Instruction Librarian	NaN	55,000	0.0	USD	NaN	NaN	United States
1	4/27/2021 11:02:22	25-34	Computing or Tech	Change & Internal Communications Manager	NaN	54,600	4000.0	GBP	NaN	NaN	United Kingdom
2	4/27/2021 11:02:38	25-34	Accounting, Banking & Finance	Marketing Specialist	NaN	34,000	NaN	USD	NaN	NaN	
3	4/27/2021 11:02:41	25-34	Nonprofits	Program Manager	NaN	62,000	3000.0	USD	NaN	NaN	
4	4/27/2021 11:02:42	25-34	Accounting, Banking & Finance	Accounting Manager	NaN	60,000	7000.0	USD	NaN	NaN	

REMOVING UNNECESSARY DATA

The research questions are not influenced by attributes like Race, Timestamp, Additional Currencies, Cities. Therefore, they were removed from the final dataset.

```
In [ ]: # Dropping unnecessary columns
init_data.drop(columns=["job_context", "compensation", "income_context", "race", "additional_currency", "timestamp"], inplace=True)
```

```
In [ ]: # Converting curr_salary column to numeric data type
init_data['curr_salary'] = init_data['curr_salary'].str.replace(",","",regex=False)
init_data['curr_salary'] = pd.to_numeric(init_data['curr_salary'])

init_data
```

```
Out[ ]: age_group curr_industry job_title curr_salary currency country state_USA city overall_exp field_exp education gender
```

0	25-34	Education (Higher Education)	Research and Instruction Librarian	55000	USD	United States	Massachusetts	Boston	5-7 years	5-7 years	Master's degree	Woman
1	25-34	Computing or Tech	Change & Internal Communications Manager	54600	GBP	United Kingdom	NaN	Cambridge	8 - 10 years	5-7 years	College degree	Non-binary
2	25-34	Accounting, Banking & Finance	Marketing Specialist	34000	USD	US	Tennessee	Chattanooga	2 - 4 years	2 - 4 years	College degree	Woman
3	25-34	Nonprofits	Program Manager	62000	USD	USA	Wisconsin	Milwaukee	8 - 10 years	5-7 years	College degree	Woman
4	25-34	Accounting, Banking & Finance	Accounting Manager	60000	USD	US	South Carolina	Greenville	8 - 10 years	5-7 years	College degree	Woman
...
27888	25-34	Computing or Tech	Systems Engineer	114000	USD	US	Wisconsin	Verona	2 - 4 years	2 - 4 years	Master's degree	Woman
27889	18-24	Business or Consulting	Data Management Consultant	60000	GBP	United Kingdom	NaN	London	2 - 4 years	2 - 4 years	College degree	Woman
27890	35-44	Engineering or Manufacturing	Project Engineer	75000	USD	United States	Missouri	Kansas City	5-7 years	5-7 years	College degree	Woman

	age_group	curr_industry	job_title	curr_salary	currency	country	state_USA	city	overall_exp	field_exp	education	gender
27891	18-24	Computing or Tech	Technology Associate	8600	EUR	Romania	NaN	Bucharest	1 year or less	1 year or less	College degree	Male
27892	35-44	Computing or Tech	Program Manager	138000	USD	USA	Texas	Houston	21 - 30 years	21 - 30 years	College degree	Woman

27893 rows x 12 columns

CONSOLIDATING ALL COUNTRY COLUMN VALUES

```
In [ ]: # Count of unique values in country column
print(len(init_data.country.unique()))
```

368

```
In [ ]: # The script below removes whitespaces from the values in the country columns, changes them to lowercase and replaces matches
countries_list = ['us', 'statesofamerica', 'america', 'usa', 'unitedstates']

init_data['country'] = init_data['country'].str.lower()
init_data['country'] = init_data['country'].str.replace('\W', '')
init_data['state_USA'] = init_data['state_USA'].replace(np.nan, 'N/A')

# checking values for every column to get top match with 70% match rate
for i in init_data.index:
    country = init_data.at[i, "country"]
    matched_country = difflib.get_close_matches(country, countries_list, n=1, cutoff=0.85)

    if matched_country:
        init_data.at[i, 'country'] = matched_country[0]

# converting values of matched countries to 'us'
for i in init_data.index:
    if init_data.at[i, "country"] in countries_list:
        init_data.at[i, "country"] = 'us'

/var/folders/h2/3d9l8ncd5xdgj1h3ffmkqyk0000gn/T/ipykernel_39323/1979729213.py:7: FutureWarning: The default value of regex
will change from True to False in a future version.
    init_data['country'] = init_data['country'].str.replace('\W', '')
```

```
In [ ]: # Checking the unique values in the country column to verify
print(len(init_data.country.unique()))
```

203

CONSOLIDATING GENDER COLUMN

```
In [ ]: # Checking NaNs in the column
print(init_data["gender"].isna().sum())

# Checking unique values for gender column
print(init_data["gender"].unique())
```

165
['Woman' 'Non-binary' 'Man' nan 'Other or prefer not to answer'
 'Prefer not to answer']

```
In [ ]: # Setting Man, Woman, and Non-Binary to M,F,NB respectively
init_data.loc[(init_data['gender']=='Man'), 'gender'] = 'M'
init_data.loc[(init_data['gender']=='Woman'), 'gender'] = 'F'
init_data.loc[(init_data['gender']=='Non-binary'), 'gender'] = 'NB'

# Setting "Other or prefer not to answer" to Other as the value "Prefer not to answer" already exists.
# Survey participants can identify as other genders that were not part of the questionnaire
init_data.loc[(init_data['gender']=='Other or prefer not to answer'), 'gender'] = 'Other'

# Setting NaN to No Answer
init_data["gender"].fillna("No Answer", inplace=True)

# Checking NaNs in the column
print(init_data["gender"].isna().sum())
```

0

CONSOLIDATING VALUES FOR CURRENT INDUSTRY

```
In [ ]: print(len(init_data["curr_industry"].unique()))
print(init_data["curr_industry"].isna().sum())
# Dropping NaN values
init_data.dropna(inplace=True)
print(init_data["curr_industry"].isna().sum())
```

1210
71
0

APPLYING FUZZY MATCHING TO BRING VALUES LIKE ["Software", "Programming", "Computing"] into IT.

```
In [ ]: # removing trailing whitespaces
init_data["curr_industry"] = init_data["curr_industry"].str.strip()

# chanign to lowercase
init_data["curr_industry"] = init_data["curr_industry"].str.lower()

# Making function for fuzzy matching
def fuzzy_match(df, col, text_to_match, text_to_replace, ratio, lim=10):

    # Getting unique values from input column
    unique_ind = df[col].unique()

    # Running fuzzy matching algorithm from fuzzywuzzy library
    matched_ind = fuzzywuzzy.process.extract(text_to_match, unique_ind, limit=lim, scorer=fuzzywuzzy.fuzz.token_sort_ratio)

    # If the similarity ratio is less than the given ratio
    if matched_ind[0][1] < ratio:
        print(f"Ratio given is too high. Try less than or equal to {matched_ind[0][1]}")
    else:
        # Create an array for best matches
        best_match = [match[0] for match in matched_ind if match[1] >= ratio]
        i = 0
        print("Matches returned :")
        # Print matches to see if the output is correct
        while best_match != [] and i < len(best_match):
            print(best_match[i])
            i += 1
        # Giving user an option to proceed or abort the matching
        ch = input("Check matches and press 'Y' to change (press and other key to abort) : ")

        # Proceeding with replacement if selection is "y"
        if ch.lower() == "y":

            # creating another column that returns boolean values for matches
            to_replace = df[col].isin(best_match)

            # replacing where boolean values are True
            df.loc[(to_replace, col)] = text_to_replace
            print(f"Replaced {text_to_match} instances in {col} column with {text_to_replace} | unique values are : {len(df)}")

    return

fuzzy_match(init_data, "curr_industry", "computing software", "IT", 59)
```

Matches returned :

- computing or tech
- software/programming
- saas company/software
- strategy consulting
- software development / it
- software development
- software products
- software
- payroll software
- biotech/software

Replaced computing software instances in curr_industry column with IT | unique values are : 983

Validating results

```
In [ ]: init_data.head(10)
```

```
Out [ ]:
```

	age_group	curr_industry	job_title	curr_salary	currency	country	state_USA	city	overall_exp	field_exp	education
0	25-34	education (higher education)	Research and Instruction Librarian	55000	USD	us	Massachusetts	Boston	5-7 years	5-7 years	Master's degree
1	25-34	IT	Change & Internal Communications Manager	54600	GBP	unitedkingdom	N/A	Cambridge	8 - 10 years	5-7 years	College degree
2	25-34	accounting, banking & finance	Marketing Specialist	34000	USD	us	Tennessee	Chattanooga	2 - 4 years	2 - 4 years	College degree
3	25-34	nonprofits	Program Manager	62000	USD	us	Wisconsin	Milwaukee	8 - 10 years	5-7 years	College degree
4	25-34	accounting, banking & finance	Accounting Manager	60000	USD	us	South Carolina	Greenville	8 - 10 years	5-7 years	College degree
5	25-34	education (higher education)	Scholarly Publishing Librarian	62000	USD	us	New Hampshire	Hanover	8 - 10 years	2 - 4 years	Master's degree
6	25-34	publishing	Publishing Assistant	33000	USD	us	South Carolina	Columbia	2 - 4 years	2 - 4 years	College degree

	age_group	curr_industry	job_title	curr_salary	currency	country	state_USA	city	overall_exp	field_exp	education
7	25-34	education (primary/secondary)	Librarian	50000	USD	us	Arizona	Yuma	5-7 years	5-7 years	Master's degree
8	45-54	IT	Systems Analyst	112000	USD	us	Missouri	St. Louis	21 - 30 years	21 - 30 years	College degree
9	35-44	accounting, banking & finance	Senior Accountant	45000	USD	us	Florida	Palm Coast	21 - 30 years	21 - 30 years	College degree



Modifying dataset so that all values are for United States for the IT industry

```
In [ ]: us_dataset = init_data.loc[(init_data['country']== 'us') & (init_data['curr_industry']=='IT')]

# Dropping cities as it is not relevant to research questions
us_dataset = us_dataset.drop('city',axis=1)
```

```
In [ ]: # Resetting index
us_dataset = us_dataset.reset_index(drop=True)
```

```
In [ ]: # Sorting values
us_dataset.sort_values(by=['curr_salary'], inplace=True)
us_dataset.head(10)
```

	age_group	curr_industry	job_title	curr_salary	currency	country	state_USA	overall_exp	field_exp	education	gender
1628	18-24	IT	Product Marketer	0	USD	us	California	1 year or less	1 year or less	College degree	M
2595	45-54	IT	Founder	0	USD	us	California	21 - 30 years	21 - 30 years	Some college	F
852	35-44	IT	Software Development Lead	1	USD	us	Wisconsin	8 - 10 years	8 - 10 years	Master's degree	M
1201	45-54	IT	Account Manager	55	USD	us	New Hampshire	21 - 30 years	8 - 10 years	College degree	F
2180	25-34	IT	Technical Writer	72	USD	us	Washington	8 - 10 years	2 - 4 years	College degree	F
2264	45-54	IT	Coach	130	USD	us	N/A	11 - 20 years	11 - 20 years	Some college	NB
1730	25-34	IT	Chief Data Scientist	240	USD	us	California	5-7 years	5-7 years	PhD	M
1038	25-34	IT	Sr Consultant	10000	USD	us	District of Columbia	8 - 10 years	8 - 10 years	Master's degree	F
3546	35-44	IT	Software Engineer Technical Support	10700	USD	us	Texas	8 - 10 years	8 - 10 years	College degree	M
819	35-44	IT	Manager of Customer Support	13000	USD	us	California	11 - 20 years	11 - 20 years	Some college	F

```
In [ ]: # Converting Job titles to lower case
us_dataset['job_title'] = us_dataset['job_title'].str.lower()
```

```
In [ ]: #software engineer dataset

# The line of code below extracts all job title swith 'software engineer'
# in the string and leaves out all job titles with senior, lead, principle, or support

se_ds = us_dataset.loc[(us_dataset['job_title'].str.contains('software engineer') & (~us_dataset['job_title'].str.contains('senior, lead, principle, or support')))
```

```
In [ ]: se_ds
```

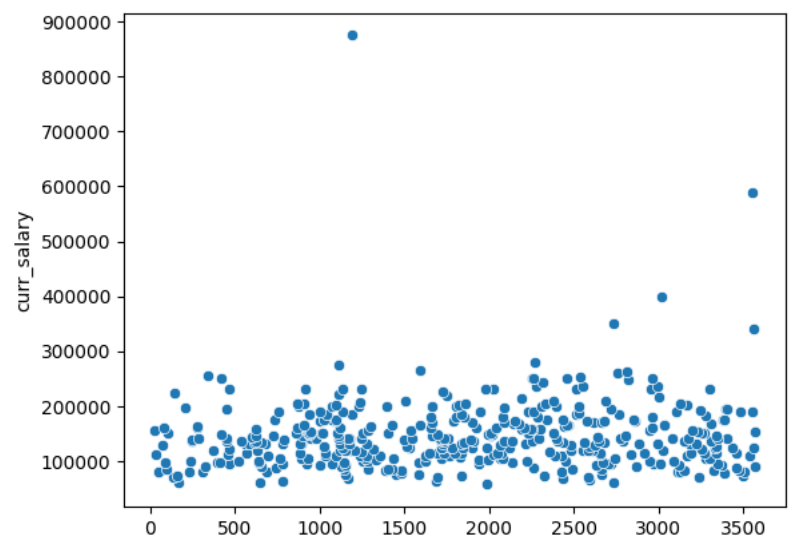
	age_group	curr_industry	job_title	curr_salary	currency	country	state_USA	overall_exp	field_exp	education	gender
1990	25-34	IT	associate software engineer	57000	USD	us	Indiana	5-7 years	1 year or less	College degree	M
167	25-34	IT	software engineer	60000	USD	us	California	5-7 years	2 - 4 years	College degree	F
649	18-24	IT	software engineer in test	60000	USD	us	New York	1 year or less	1 year or less	College degree	F
2731	25-34	IT	software engineer	61000	USD	us	California	5-7 years	1 year or less	College degree	F
1690	35-44	IT	developer (software engineer/programmer)	63500	USD	us	Missouri	11 - 20 years	2 - 4 years	Some college	M
...
3561	25-34	IT	software engineer	340000	USD	us	Michigan, Texas, Washington	2 - 4 years	2 - 4 years	Master's degree	M
2738	25-34	IT	software engineer	350000	USD	us	California	2 - 4 years	2 - 4	College	F

	age_group	curr_industry	job_title	curr_salary	currency	country	state_USA	overall_exp	field_exp	education	gender
									years	degree	
3014	25-34	IT	software engineer	400000	USD	us	Massachusetts	5-7 years	5-7 years	PhD	M
3551	25-34	IT	software engineer	590000	USD	us	New York	8 - 10 years	8 - 10 years	College degree	F
1191	25-34	IT	software engineer	875000	USD	us	California	5-7 years	2 - 4 years	College degree	M

406 rows x 11 columns

Searching for outliers

```
In [ ]: # Creating scatter plot for outliers
sns.scatterplot(y=se_ds['curr_salary'],x=se_ds.index)
plt.show()
```



Most of the vlaues are concentrated around 100,000 and around 250,000

```
In [ ]: # Removing outliers from the dataset using Inter-Quartile Range (IQR)
# IQR depicts the spread of values in the current salary column

qt_1, qt_3 = np.percentile(se_ds['curr_salary'],[25,75])
iqr = qt_3-qt_1
lower_bound = qt_1 - (1.5 * iqr)
upper_bound = qt_3 + (1.5 * iqr)

print(lower_bound, upper_bound)

se_ds = se_ds.loc[(se_ds['curr_salary'] > lower_bound) & (se_ds['curr_salary'] < upper_bound)]
```

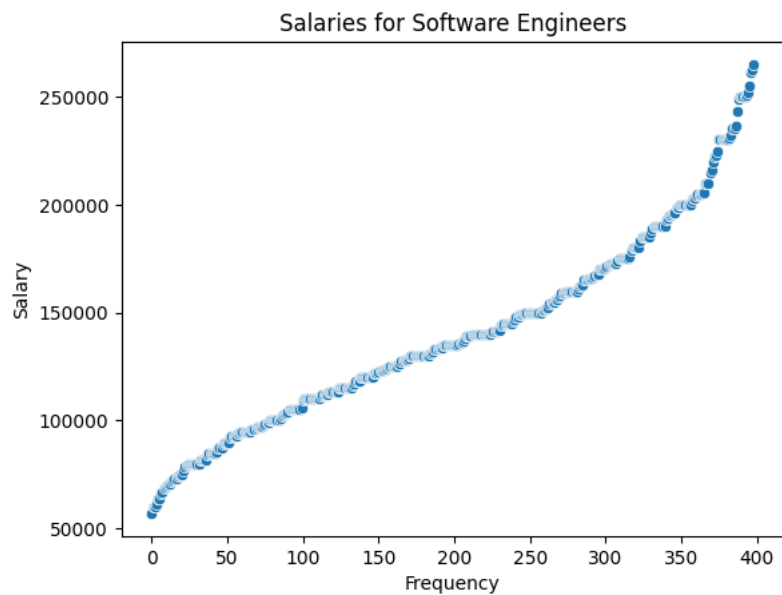
16250.0 266250.0

```
In [ ]: # resetting index
se_ds = se_ds.reset_index(drop=True)
```

```
In [ ]: # Describing dataset
se_ds.describe()
```

```
Out[ ]:      curr_salary
count      399.000000
mean      140925.408521
std        45374.261953
min        57000.000000
25%       107725.000000
50%       135000.000000
75%       170200.000000
max       265000.000000
```

```
In [ ]: sns.scatterplot(y=se_ds['curr_salary'],x=se_ds.index)
plt.title("Salaries for Software Engineers")
plt.xlabel('Frequency')
plt.ylabel('Salary')
plt.show()
```



```
In [ ]: # Exporting to Excel for Visualization in Power BI

filename = 'se_dataset.xlsx'
se_ds.to_excel(filename)
```

```
In [ ]:
```

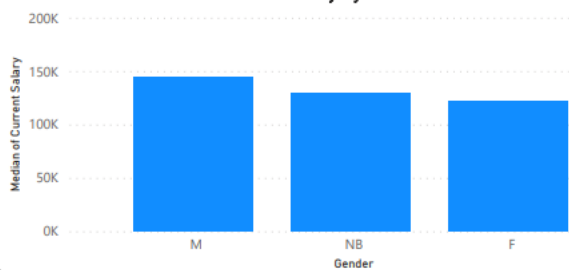
Software Engineers in US.pdf

1 / 1

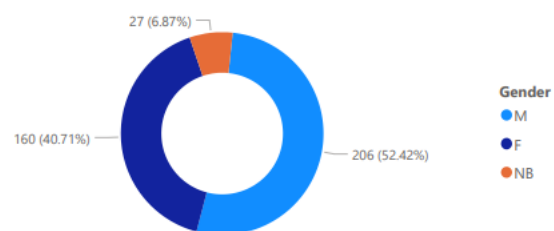
42%



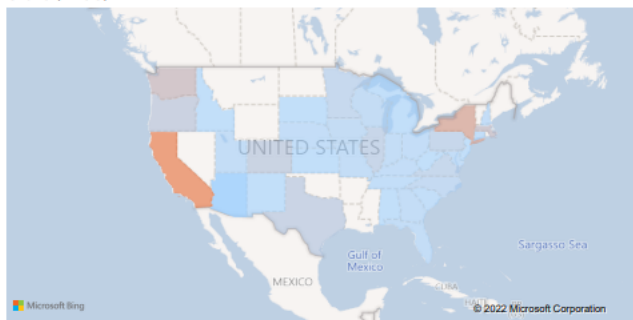
Median Salary by Gender



Gender Composition



State (in US)



Median of Current Salary by Education and Gender

